

# Attention Is All You Need

- Authors: *Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, Illia Polosukhin*
  - Main Affiliation: Google Brain
- 

2025.07.22

서정호

# Contents

---

- 1. Proposal**
- 2. Model Architecture**
- 3. Why Self-Attention**

# 1. Proposal

## NLP → Sequence transduction model.

- Sequence → Sequence 변환
- 기존 (1): RNN or CNN 사용. using encoder & decoder
- 기존 (2): Best model: Attention mechanism 사용해서 encoder - decoder 연결

## Transformer architecture 제안.

- RNN[LSTM, GRU], CNN 사용 X
- 오직 attention mechanisms [specifically, self-attention]만 사용.
- 기존: encoder-decoder architectures를 사용한 recurrent layers  
→ transformer: multi-headed self-attention
- Experiments 결과 → transformer: 병렬화 up & 학습시간 down.
- 평가척도: BLEU score, 2014 WMT [Workshop on Machine Translation]에서 성과 얻음.
  - English to German, English to French
- GPU도 적게 사용!!.

## 2. Model Architecture

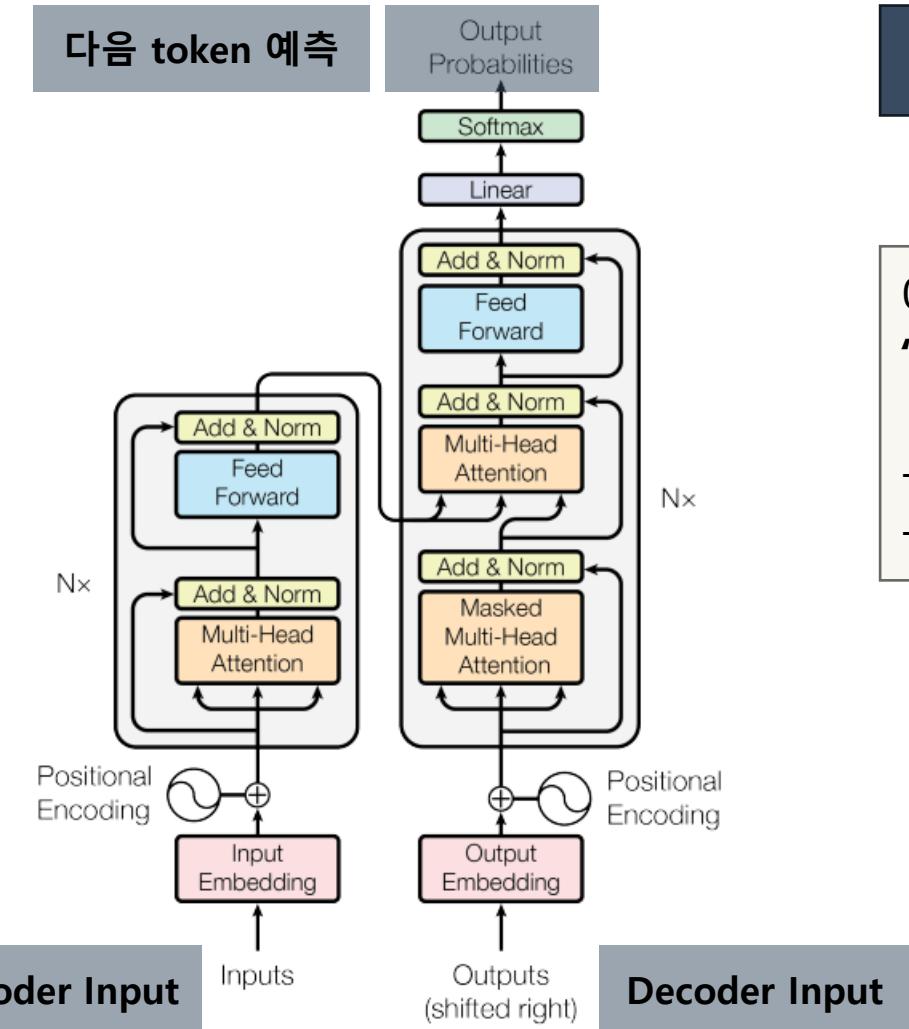
---

### ▼ Sequence transduction models

#### encoder-decoder structure

- Encoder: [input]  $\mathbf{x} \rightarrow \mathbf{z}$  [output]
  - $\mathbf{x} = (x_1, \dots, x_n)$ : Symbol sequence
  - $\mathbf{z} = (z_1, \dots, z_n)$ : Continuous sequence
- Decoder: [input]  $\mathbf{z} \rightarrow \mathbf{y}$  [output]
  - $\mathbf{z} = (z_1, \dots, z_n)$ : Continuous sequence
  - $\mathbf{y} = (y_1, \dots, y_m)$ : Symbol sequence

## 2. Model Architecture



### Total Architecture

예. Eng. -> French.  
“I am a student” -> “Je suis étudiant »

- Input token: [I, am, a, student]
- Target token: [Je, suis, étudiant]

Figure 1: The Transformer - model architecture.

## 2. Model Architecture

### Learning

예. Eng. -> French.

"I am a student" -> "Je suis étudiant »

- Input token: [I, am, a, student]
- Target token: [Je, suis, étudiant]

### 입력

- **Encoder input:** [I, am, a, student]
- **Decoder input** (shifted right): [<s>, Je, suis]
- **Target** (예측 정답): [Je, suis, étudiant]

### 학습 목표

각 시점  $t$ 에서 다음을 예측하도록 학습:

maximized

Step $t$	Decoder Input	Target Token	모델이 학습할 조건부 확률	Loss ftn.
1	시작 token	[<s>]	Je	$P(\text{Je}   < s >, x)$
2		suis	$P(\text{suis}   < s >, \text{Je}, x)$	
3		étudiant	$P(\text{étudiant}   < s >, \text{Je}, \text{suis}, x)$	

## 2. Model Architecture

### Inference

예. Eng. -> French.  
**"I am a student"** -> **"Je suis étudiant"**

- Input token: **[I, am, a, student]**
- Target token: **[Je, suis, étudiant]**

### 입력

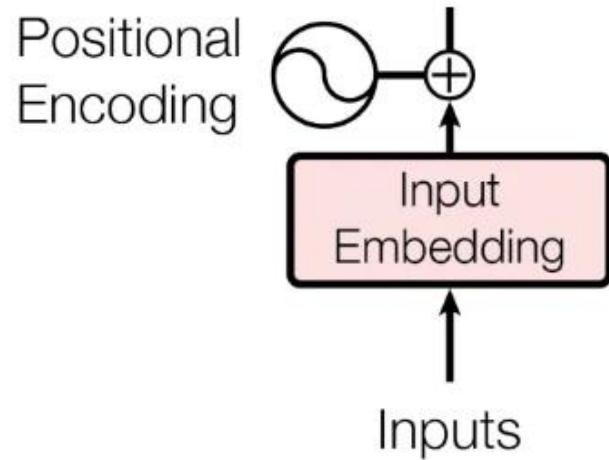
- **Encoder input:** **[I, am, a, student]**
- **Decoder 시작:** **[<s>]** → 이후 토큰은 모델이 생성

### 디코딩 과정 (Greedy 예시)

Step $t$	Decoder Input	Softmax Output (예시)	선택된 토큰
1	<b>[&lt;s&gt;]</b>	{"Je": 0.8, "Tu": 0.1, ...}	<b>Je</b>
2	<b>[&lt;s&gt;, Je]</b>	{"suis": 0.9, "vais": 0.05, ...}	<b>suis</b>
3	<b>[&lt;s&gt;, Je, suis]</b>	{"étudiant": 0.95, ...}	<b>étudiant</b>
4	<b>[&lt;s&gt;, Je, suis, étudiant]</b>	{"</s>": 0.99, ...}	<b>&lt;/s&gt;</b> → 종료 <b>종료 token</b>

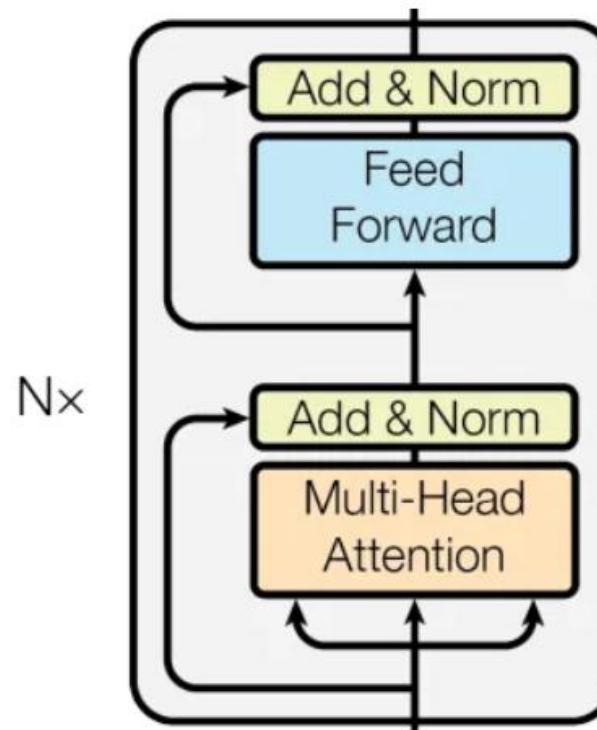
## 2. Model Architecture

▼ 1st.



- Inputs == input tokens
- tokens → embedding vector  $\in \mathbb{R}^{d_{\text{model}}}$
- embedding vector + positional encoding: Input.

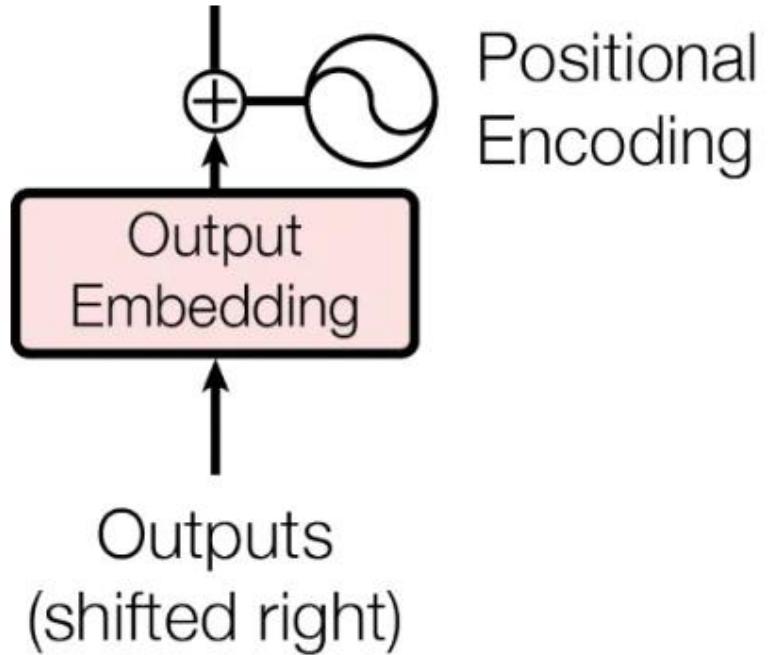
▼ 2nd. == Encoder



- $N = 6$  개의 identical layers, input → 1st layer → 2nd layer → ... →  $N$ -th layer → output
- each layer == two sub layers
  1. [Multi-Head Attention] Multi-Head self-Attention mechanism
  2. [Feed Forward] position-wise fully connected Feed-Forward Network [FFN]
    - [Add & Norm] LayerNorm () function
- Sub-layers' outputs
  - LayerNorm ( $\mathbf{x} + \text{Sublayer}(\mathbf{x})$ ), where  $\mathbf{x}$ : input of the each sub-layer

## 2. Model Architecture

### ▼ 3rd.

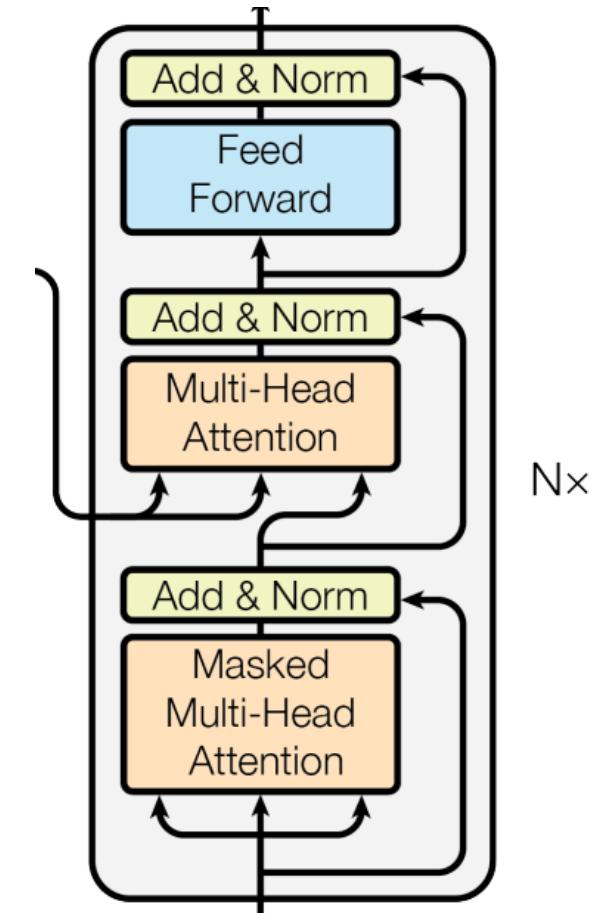


- Outputs == output tokens
- tokens → embedding vector  $\in \mathbb{R}^{d_{\text{model}}}$
- embedding vector + positional encoding: Output.

## 2. Model Architecture

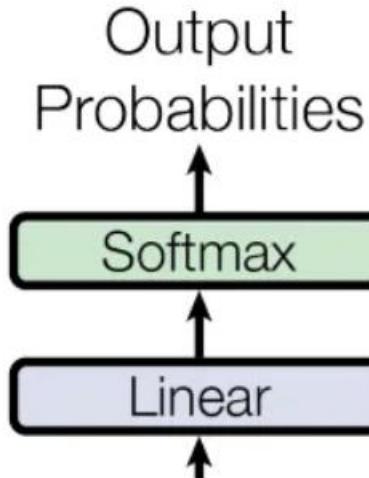
### ▼ 4th. == Decoder

- N = 6개의 identical layers, input → 1st layer → 2nd layer → ... → N-th layer → output
- each layer == three sub layers
  1. [Masked Multi-Head Attention]
    - Masking: i의 position이 i보다 작은 positions의 outputs에만 의존함을 보장.
    - 1, 2, 3, ..., i-1 → i
  2. [Multi-Head Attention] Multi-Head self-Attention mechanism
    - 주의: Encoder output에 multi-head attention 적용.
  3. [Feed Forward] position-wise fully connected Feed-Forward Network [FFN]
    - [Add & Norm] LayerNorm () function
- Sub-layers' outputs
  - LayerNorm ( $\mathbf{x} + \text{Sublayer}(\mathbf{x})$ ), where  $\mathbf{x}$ : input of the each sub-layer



## 2. Model Architecture

▼ 5th. = 3.4.



# Total Architecture

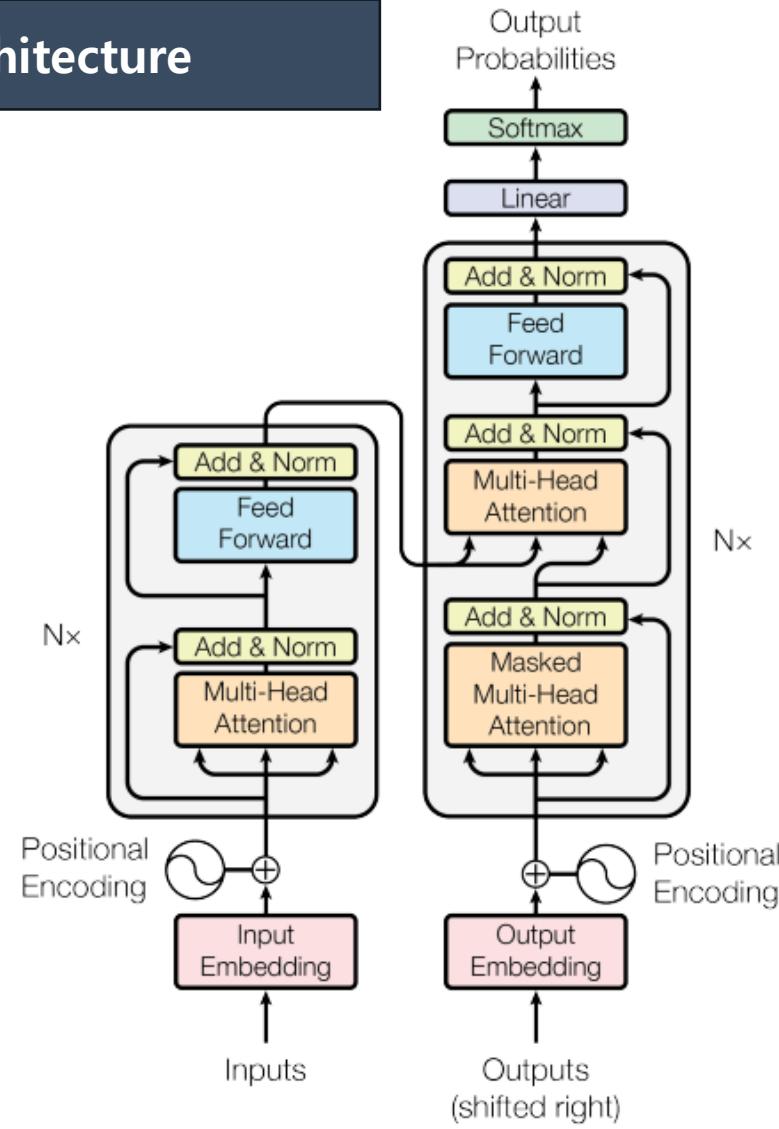


Figure 1: The Transformer - model architecture.

# 2. Model Architecture

## ▼ 3.2. Attention

- Inputs
  - query vector:  $\mathbf{q} \rightarrow$  dimension  $d_q = d_k$
  - keys vector:  $\mathbf{k} \rightarrow$  dimension  $d_k$
  - values vector:  $\mathbf{v} \rightarrow$  dimension  $d_v$
- Output vector = values의 가중합, weight  $\leftarrow$  (query, key)

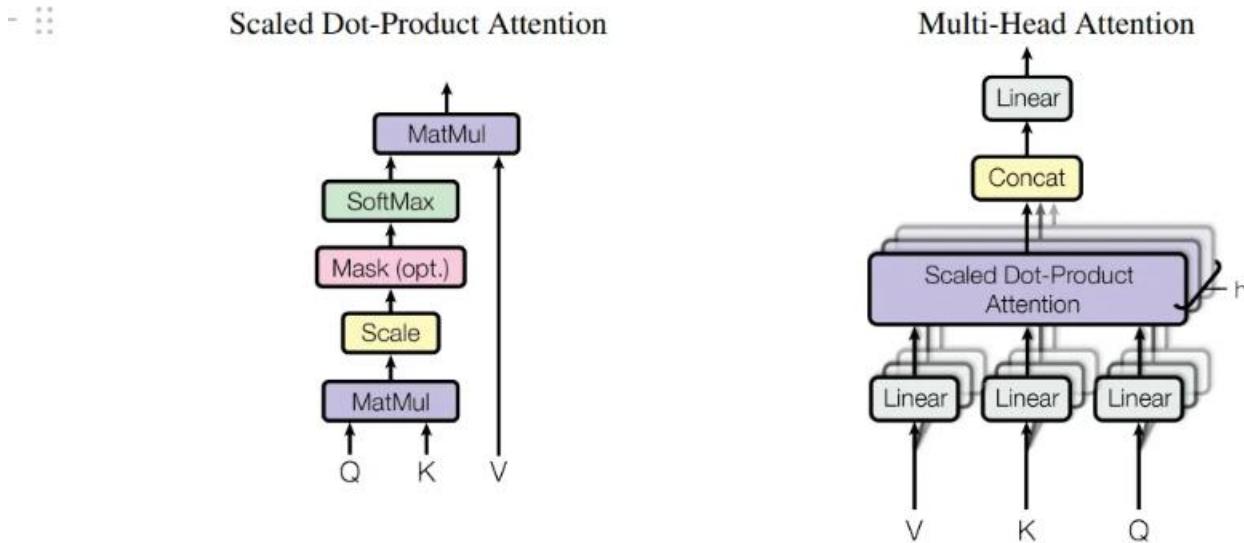
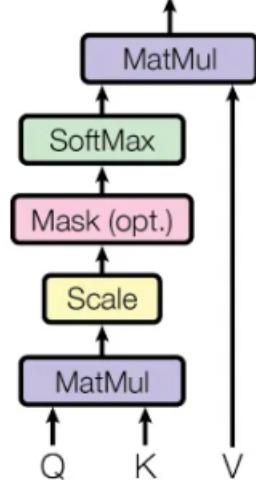


Figure 2: (left) Scaled Dot-Product Attention. (right) Multi-Head Attention consists of several attention layers running in parallel.

## 2. Model Architecture

### ▼ 3.2.1. Scaled Dot-Product Attention [single attention head]

Scaled Dot-Product Attention

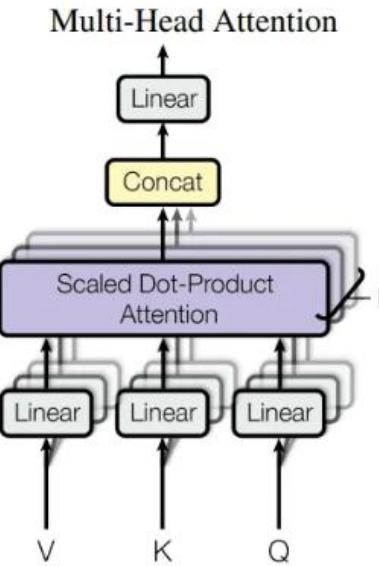


- $\mathbf{q}, \mathbf{k}, \mathbf{v} \rightarrow \text{matrix } Q, K, V$

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \quad (1)$$

- dot product → scaling → softmax

### ▼ 3.2.2. Multi-Head Attention



- $h = 8$ 개의 multi-head 사용. with  $W$  : parameter matrix

$$\begin{aligned} \text{MultiHead}(Q, K, V) &= \text{Concat}(\text{head}_1, \dots, \text{head}_h) W^O \\ \text{where } \text{head}_i &= \text{Attention}\left(QW_i^Q, KW_i^K, VW_i^V\right) \end{aligned}$$

$$W_i^Q \in \mathbb{R}^{d_{\text{model}} \times d_k}, W_i^K \in \mathbb{R}^{d_{\text{model}} \times d_k}, W_i^V \in \mathbb{R}^{d_{\text{model}} \times d_v}$$

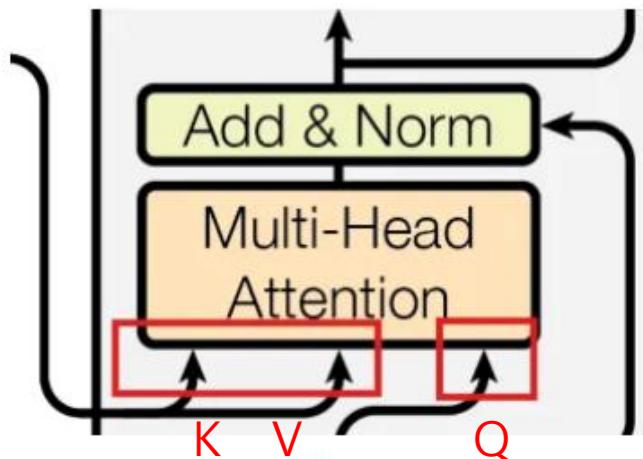
$$d_k = d_v = \frac{d_{\text{model}}}{h} = 64 \rightarrow d_{\text{model}} = 512$$

## 2. Model Architecture

### ▼ 3.2.3. Applications of Attention in our Model

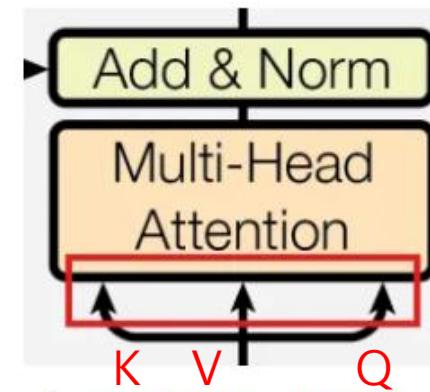
Transformer uses multi-head attention in three different ways

1. Encoder-Decoder Attention layers



- a. previous decoder layer's output → queries
- b. encoder layer's output → memory keys and values

2. Encoder's self-attention layers [Q, K, V가 모두 같은 곳에서 나옴.]



1. previous encoder layer's output → keys, values and queries
3. Decoder's self-attention layers

## 2. Model Architecture

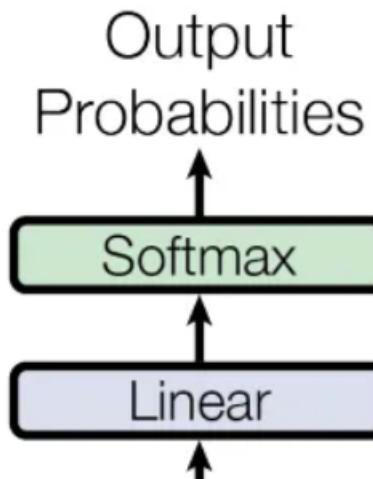
### 3.3. Position-wise Feed-Forward Networks

- [Feed Forward]: Fully connected feed-forward network

$$\text{FFN}(\mathbf{x}) = \text{RELU}(\mathbf{x}W_1 + \mathbf{b}_1)W_2 + \mathbf{b}_2 \quad (2)$$

where  $\text{RELU}(x) = \max(0, x)$

### ▼ 3.4. Embeddings and Softmax



- input/output tokens → embedding vectors
- Learned linear transformation & softmax function: decoder output → predicted next-token probabilities.

## 2. Model Architecture

### Total Architecture

#### ▼ 3.5. Positional Encoding

- RNN, CNN 사용 X → sequence의 ordering 사용 불가.
  - → Positional encodings를 encoder and decoder stacks의 bottoms에 삽입.
- Sequence의 ordering property 사용 위함.
- 다양한 positional encodings 사용 가능. → 논문에서는 아래와 같이 사용.

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right)$$

$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right)$$

- pos: position
- i: dimension

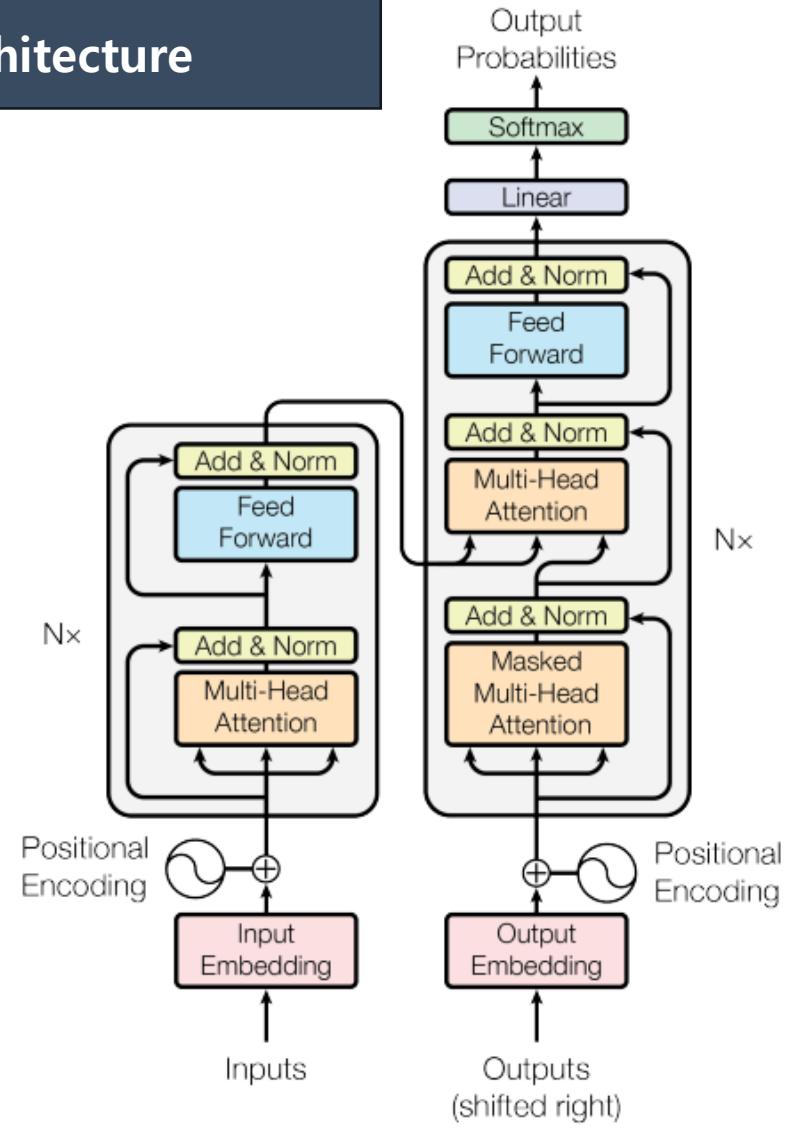


Figure 1: The Transformer - model architecture.

### 3. Why Self-Attention

Table 1: Maximum path lengths, per-layer complexity and minimum number of sequential operations for different layer types.  $n$  is the sequence length,  $d$  is the representation dimension,  $k$  is the kernel size of convolutions and  $r$  the size of the neighborhood in restricted self-attention.

Layer Type	Complexity per Layer	Sequential Operations	Maximum Path Length
Self-Attention	$O(n^2 \cdot d)$	$O(1)$	$O(1)$
Recurrent	$O(n \cdot d^2)$	$O(n)$	$O(n)$
Convolutional	$O(k \cdot n \cdot d^2)$	$O(1)$	$O(\log_k(n))$
Self-Attention (restricted)	$O(r \cdot n \cdot d)$	$O(1)$	$O(n/r)$